# 趋境科技 8 卡 H200 方案 DeepSeek-V3.2 测试报告

**测试时间**：2026.01.10

## 1. 测试环境

### 1.1 硬件配置

| 硬件名称 | 配置信息 | 数量 |
| --- | --- | --- |
| CPU | Intel Xeon Platinum 8558 | 2 |
| GPU | NVIDIA H200 | 8 |
| 内存 | DDR5 64GB | 32 条 |

**服务器数量**：3 台

### 1.2 软件配置

| 软件名称 | 版本信息 |
| --- | --- |
| 操作系统 | Ubuntu 22.04.5 LTS |
| 内核版本 | 5.15.0-164-generic |
| 测试工具 | evalscope v1.4.1 |

# 2. 测试用例

## 2.1 性能测试

**测试目的**：验证趋境方案 DeepSeek-V3.2 性能表现（TTFT/TPS 等）符合预期。

**测试命令**：

```
evalscope perf --rate 5 --parallel 84 --number 512 \
  --model DeepSeek-V3.2 \
  --tokenizer-path /data/models/DeepSeek-V3.2 \
  --url http://127.0.0.1:8000/v1/completions \
  --api openai --dataset random \
  --max-tokens 1000 --min-tokens 1000 \
  --prefix-length 0 --min-prompt-length 6000 --max-prompt-length 6000 \
  --extra-args '{"ignore_eos": true}'
```

**预期结果**：TTFT ⩽ 2s；TPOT ⩽ 20ms

**测试结果**：

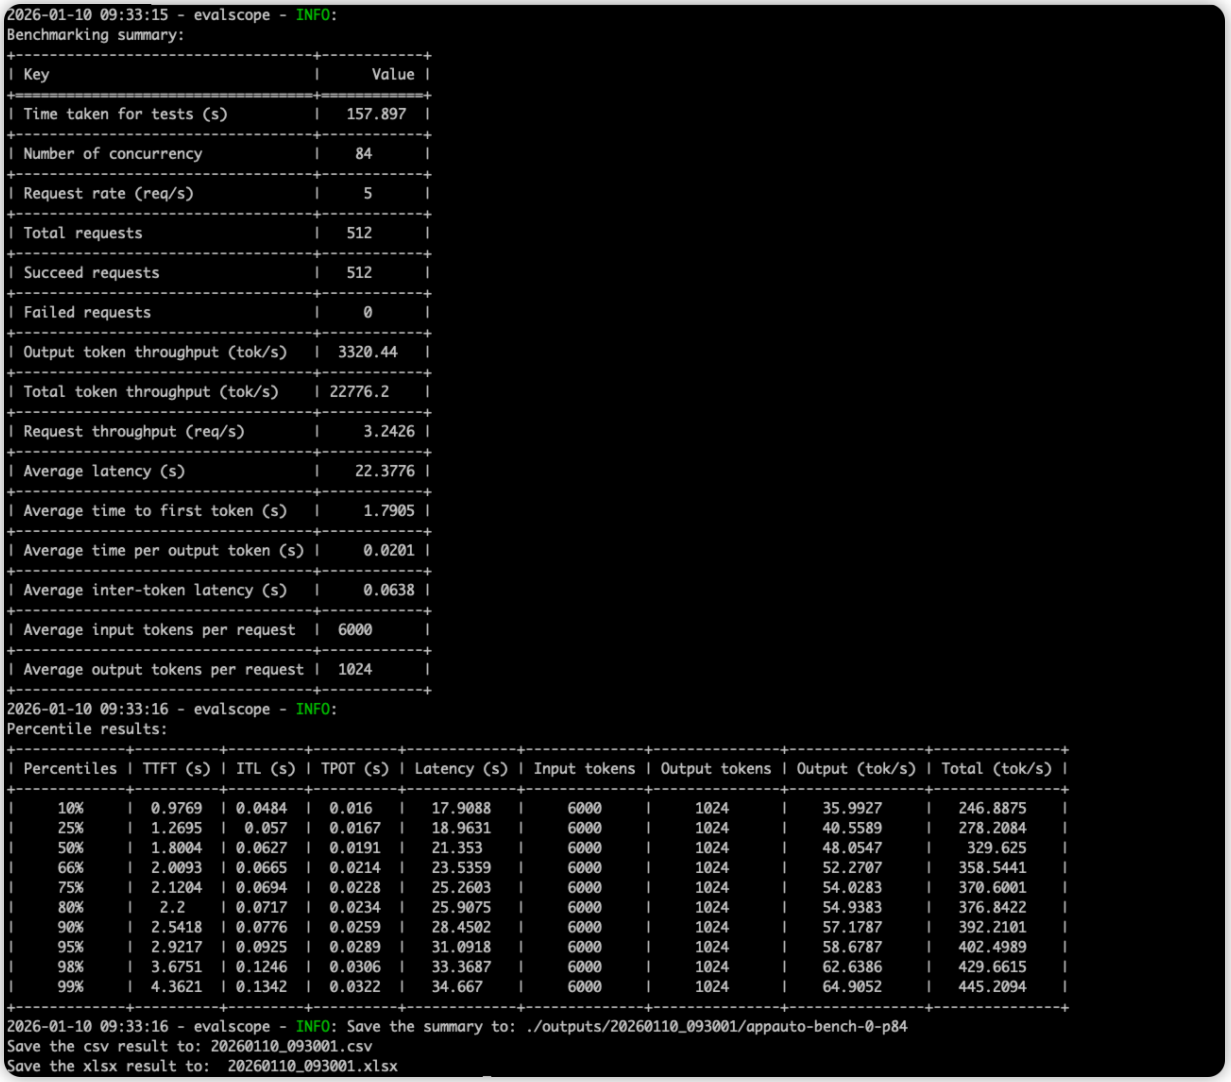

```
2026-01-10 09:33:15 - evalscope - INFO:
Benchmarking summary:
+---------------------------------+------------+
| Key                             |      Value |
+=================================+============+
| Time taken for tests (s)        |   157.897  |
+---------------------------------+------------+
| Number of concurrency           |    84      |
+---------------------------------+------------+
| Request rate (req/s)            |     5      |
+---------------------------------+------------+
| Total requests                  |   512      |
+---------------------------------+------------+
| Succeed requests                |   512      |
+---------------------------------+------------+
| Failed requests                 |     0      |
+---------------------------------+------------+
| Output token throughput (tok/s) |  3320.44   |
+---------------------------------+------------+
| Total token throughput (tok/s)  | 22776.2    |
+---------------------------------+------------+
| Request throughput (req/s)      |     3.2426 |
+---------------------------------+------------+
| Average latency (s)             |    22.3776 |
+---------------------------------+------------+
| Average time to first token (s) |     1.7905 |
+---------------------------------+------------+
| Average time per output token (s) |     0.0201 |
+---------------------------------+------------+
| Average inter-token latency (s) |     0.0638 |
+---------------------------------+------------+
| Average input tokens per request |  6000      |
+---------------------------------+------------+
| Average output tokens per request |  1024      |
+---------------------------------+------------+
2026-01-10 09:33:16 - evalscope - INFO:
Percentile results:
+-------------+----------+---------+----------+-------------+--------------+---------------+----------------+---------------+
| Percentiles | TTFT (s) | ITL (s) | TPOT (s) | Latency (s) | Input tokens | Output tokens | Output (tok/s) | Total (tok/s) |
+-------------+----------+---------+----------+-------------+--------------+---------------+----------------+---------------+
|     10%     |  0.9769  | 0.0484  |  0.016   |   17.9088   |     6000     |     1024      |    35.9927     |   246.8875    |
|     25%     |  1.2695  |  0.057  |  0.0167  |   18.9631   |     6000     |     1024      |    40.5589     |   278.2084    |
|     50%     |  1.8004  | 0.0627  |  0.0191  |   21.353    |     6000     |     1024      |    48.0547     |    329.625    |
|     66%     |  2.0093  | 0.0665  |  0.0214  |   23.5359   |     6000     |     1024      |    52.2707     |   358.5441    |
|     75%     |  2.1204  | 0.0694  |  0.0228  |   25.2603   |     6000     |     1024      |    54.0283     |   370.6001    |
|     80%     |   2.2    | 0.0717  |  0.0234  |   25.9075   |     6000     |     1024      |    54.9383     |   376.8422    |
|     90%     |  2.5418  | 0.0776  |  0.0259  |   28.4502   |     6000     |     1024      |    57.1787     |   392.2101    |
|     95%     |  2.9217  | 0.0925  |  0.0289  |   31.0918   |     6000     |     1024      |    58.6787     |   402.4989    |
|     98%     |  3.6751  | 0.1246  |  0.0306  |   33.3687   |     6000     |     1024      |    62.6386     |   429.6615    |
|     99%     |  4.3621  | 0.1342  |  0.0322  |   34.667    |     6000     |     1024      |    64.9052     |   445.2094    |
+-------------+----------+---------+----------+-------------+--------------+---------------+----------------+---------------+
2026-01-10 09:33:16 - evalscope - INFO: Save the summary to: ./outputs/20260110_093001/appauto-bench-0-p84
Save the csv result to: 20260110_093001.csv
Save the xlsx result to:  20260110_093001.xlsx
```

**结论**：✅ 通过

---

## 2.2 正确性测试

**测试目的**：验证趋境方案 DeepSeek-V3.2 精度表现符合预期。

**测试命令**：

```
evalscope eval --model DeepSeek-V3.2 \
  --api-url http://127.0.0.1:8000/v1/ --api-key EMPTY \
  --eval-type openai_api --datasets 'aime25' \
  --stream --timeout 6000 \
  --generation-config '{"do_sample":true,"temperature":0.6,"max_tokens":56000,"ex
```

```
  --dataset-args '{"aime25": {"prompt_template": "{question}\nPlease reason step
  --eval-batch-size 16
```

**评分标准**：官方分数

**测试结果**：

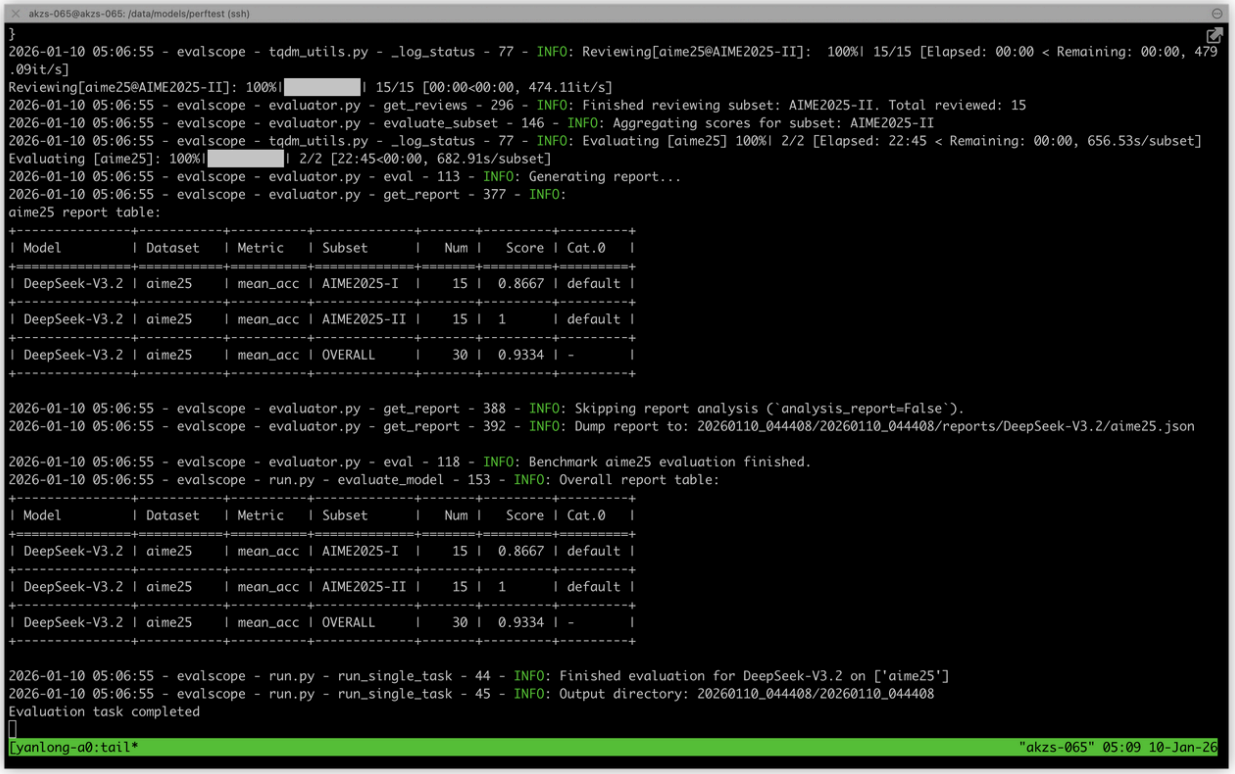


**结论**：✅ 通过（趋境：93.3；官方：93.1）

---

## 2.3 超长上下文测试

**测试目的**：验证趋境方案 DeepSeek-V3.2 支持超长上下文（128K）。

**测试命令**：

```
evalscope perf --parallel 1 --number 1 \
  --model DeepSeek-V3.2 \
  --tokenizer-path /data/models/DeepSeek-V3.2 \
  --url http://127.0.0.1:8000/v1/completions \
  --api openai --dataset random \
  --max-tokens 1000 --min-tokens 1000 \
```

```
  --prefix-length 0 --min-prompt-length 128000 --max-prompt-length 128000 \
  --extra-args '{"ignore_eos": true}'
```

**预期结果**：支持 128K 超长上下文

**测试结果**：

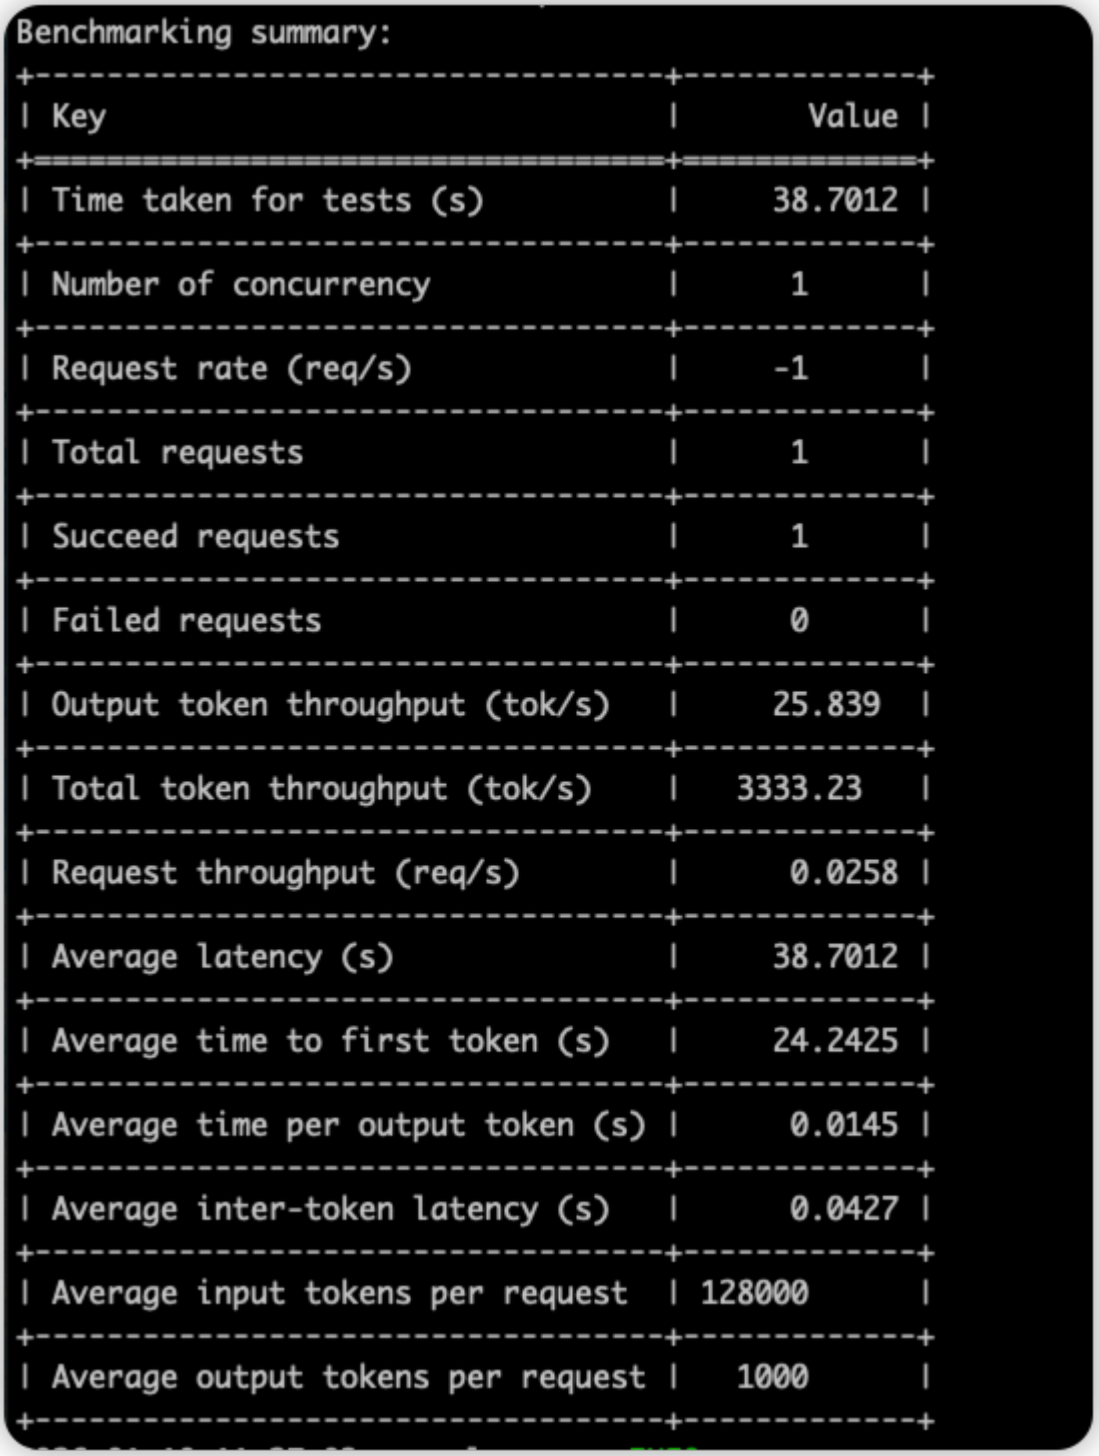

```
Benchmarking summary:
+-----------------------------------+-------------+
| Key                               |       Value |
+===================================+=============+
| Time taken for tests (s)          |     38.7012 |
+-----------------------------------+-------------+
| Number of concurrency             |      1      |
+-----------------------------------+-------------+
| Request rate (req/s)              |     -1      |
+-----------------------------------+-------------+
| Total requests                    |      1      |
+-----------------------------------+-------------+
| Succeed requests                  |      1      |
+-----------------------------------+-------------+
| Failed requests                   |      0      |
+-----------------------------------+-------------+
| Output token throughput (tok/s)   |     25.839  |
+-----------------------------------+-------------+
| Total token throughput (tok/s)    |   3333.23   |
+-----------------------------------+-------------+
| Request throughput (req/s)        |      0.0258 |
+-----------------------------------+-------------+
| Average latency (s)               |     38.7012 |
+-----------------------------------+-------------+
| Average time to first token (s)   |     24.2425 |
+-----------------------------------+-------------+
| Average time per output token (s) |      0.0145 |
+-----------------------------------+-------------+
| Average inter-token latency (s)   |      0.0427 |
+-----------------------------------+-------------+
| Average input tokens per request  | 128000      |
+-----------------------------------+-------------+
| Average output tokens per request |   1000      |
+-----------------------------------+-------------+
```

**结论**：✅ 通过

---

## 3. 测试结论

三台服务器的整体 TPM 为 **1,366,560**（22,776 × 60），TPOT、TTFT、精度均符合要求。

结合上述测试结果，**测试通过**。当前性能仍在持续优化中。